

titanic

July 6, 2024

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import math
%matplotlib inline

titanic_data = pd.read_csv("Titanic.csv")
titanic_data.head(10)
```

```
[1]: PassengerId  Survived  Pclass    Sex    Age  SibSp  Parch    Fare  \
0             1         0        3   Male   22.0     1     0    7.2500
1             2         1        1  female  38.0     1     0   71.2833
2             3         1        3  female  26.0     0     0    7.9250
3             4         1        1  female  35.0     1     0   53.1000
4             5         0        3   Male   35.0     0     0    8.0500
5             6         0        3   Male   60.0     0     0    8.4583
6             7         0        1   Male   54.0     0     0   51.8625
7             8         0        3   Male    2.0     3     1   21.0750
8             9         1        3  female  27.0     0     2   11.1333
9            10         1        2  female  14.0     1     0   30.0708
```


```
Embarked
0        3
1        1
2        3
3        3
4        3
5        2
6        3
7        3
8        3
9        1
```

```
[2]: print('# of passengers in original data:' +str(len(titanic_data.index)))
```

```
# of passengers in original data:889
```


```
[3]: sns.countplot(x="Survived", data=titanic_data)
```

```
[3]: <Axes: xlabel='Survived', ylabel='count'>
```




```
[4]: sns.countplot(x="Survived", hue="Sex", data=titanic_data)
```

```
[4]: <Axes: xlabel='Survived', ylabel='count'>
```




```
[5]: sns.countplot(x="Survived",hue="Pclass",data=titanic_data)
```

```
[5]: <Axes: xlabel='Survived', ylabel='count'>
```




```
[6]: titanic_data['Age'].plot.hist()
```

```
[6]: <Axes: ylabel='Frequency'>
```




```
[7]: titanic_data['Fare'].plot.hist()
```

```
[7]: <Axes: ylabel='Frequency'>
```




```
[55]: titanic_data['Fare'].plot.hist(bins=20, figsize=(10,5))
```

```
[55]: <Axes: ylabel='Frequency'>
```



```
[10]: titanic_data['Fare'].plot.hist(bins=20, figsize=(10,5))
```

```
[10]: <Axes: ylabel='Frequency'>
```




```
[11]: titanic_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 889 entries, 0 to 888
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId 889 non-null   int64
1   Survived    889 non-null   int64
2   Pclass      889 non-null   int64
3   Sex         889 non-null   object
4   Age         889 non-null   float64
5   SibSp       889 non-null   int64
6   Parch       889 non-null   int64
7   Fare        889 non-null   float64
8   Embarked    889 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 62.6+ KB
```

```
[12]: sns.countplot(x='SibSp', data=titanic_data)
```

```
[12]: <Axes: xlabel='SibSp', ylabel='count'>
```



```
[13]: titanic_data.isnull()
```

```
[13]:
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	\
0	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	
..	
884	False	False	False	False	False	False	False	False	
885	False	False	False	False	False	False	False	False	
886	False	False	False	False	False	False	False	False	
887	False	False	False	False	False	False	False	False	
888	False	False	False	False	False	False	False	False	

	Embarked
0	False
1	False
2	False
3	False
4	False


```
..      ...
884      False
885      False
886      False
887      False
888      False
```


```
[889 rows x 9 columns]
```

```
[14]: titanic_data.isnull().sum()
```

```
[14]: PassengerId      0
      Survived        0
      Pclass          0
      Sex             0
      Age             0
      SibSp           0
      Parch           0
      Fare            0
      Embarked        0
      dtype: int64
```

```
[15]: sns.boxplot(x='Pclass', y='Age', data=titanic_data)
```

```
[15]: <Axes: xlabel='Pclass', ylabel='Age'>
```



```
[16]: titanic_data.head(5)
```

```
[16]:
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	\
0	1	0	3	Male	22.0	1	0	7.2500	
1	2	1	1	female	38.0	1	0	71.2833	
2	3	1	3	female	26.0	0	0	7.9250	
3	4	1	1	female	35.0	1	0	53.1000	
4	5	0	3	Male	35.0	0	0	8.0500	

	Embarked
0	3
1	1
2	3
3	3
4	3

```
[17]: titanic_data.dropna(inplace=True)
```

```
[18]: titanic_data.isnull().sum()
```

```
[18]: PassengerId    0
      Survived      0
      Pclass        0
      Sex           0
      Age           0
      SibSp         0
      Parch         0
      Fare          0
      Embarked      0
      dtype: int64
```

```
[19]: titanic_data.head(2)
```

```
[19]:   PassengerId  Survived  Pclass    Sex  Age  SibSp  Parch    Fare  \
0             1         0       3   Male  22.0     1     0    7.2500
1             2         1       1  female  38.0     1     0   71.2833

      Embarked
0           3
1           1
```

```
[20]: pd.get_dummies(titanic_data['Sex'])
```

```
[20]:   Male  female
0    True   False
1   False    True
2   False    True
3   False    True
4    True   False
..    ...    ...
884   True   False
885  False    True
886  False    True
887   True   False
888   True   False

[889 rows x 2 columns]
```

```
[22]: sex = pd.get_dummies(titanic_data['Sex'], drop_first=True)
      sex.head(5)
```

```
[22]:   female
0   False
1    True
2    True
3    True
4   False
```

```
[23]: embark=pd.get_dummies(titanic_data["Embarked"],drop_first=True)
      embark.head(5)
```

```
[23]:      2      3
0  False   True
1  False  False
2  False   True
3  False   True
4  False   True
```

```
[24]: Pcl=pd.get_dummies(titanic_data['Pclass'],drop_first=True)
      Pcl.head(5)
```

```
[24]:      2      3
0  False   True
1  False  False
2  False   True
3  False  False
4  False   True
```

```
[25]: titanic_data=pd.concat([titanic_data,sex,embark,Pcl],axis=1)
```

```
[26]: titanic_data.head(5)
```

```
[26]:   PassengerId  Survived  Pclass    Sex  Age  SibSp  Parch    Fare  \
0             1         0        3   Male  22.0     1     0   7.2500
1             2         1        1  female  38.0     1     0  71.2833
2             3         1        3  female  26.0     0     0   7.9250
3             4         1        1  female  35.0     1     0  53.1000
4             5         0        3   Male  35.0     0     0   8.0500

      Embarked  female      2      3      2      3
0             3  False  False   True  False   True
1             1   True  False  False  False  False
2             3   True  False   True  False   True
3             3   True  False   True  False  False
4             3  False  False   True  False   True
```

```
[ ]: titanic_data.drop(['Sex','Embarked','PassengerId'],axis=1,inplace=True)
```

```
[28]: titanic_data.head()
```

```
[28]:   PassengerId  Survived  Pclass    Sex  Age  SibSp  Parch    Fare  \
0             1         0        3   Male  22.0     1     0   7.2500
1             2         1        1  female  38.0     1     0  71.2833
2             3         1        3  female  26.0     0     0   7.9250
3             4         1        1  female  35.0     1     0  53.1000
```

```
4          5          0          3    Male  35.0          0          0    8.0500
```

```
    Embarked  female          2          3          2          3
0          3    False  False    True  False    True
1          1     True  False  False  False  False
2          3     True  False    True  False    True
3          3     True  False    True  False  False
4          3    False  False    True  False    True
```

```
[29]: titanic_data.drop(['Pclass'],axis=1,inplace=True)
```

```
[30]: titanic_data.head()
```

```
[30]: PassengerId  Survived  Sex   Age  SibSp  Parch    Fare  Embarked  \
0             1         0   Male  22.0     1     0   7.2500         3
1             2         1 female  38.0     1     0  71.2833         1
2             3         1 female  26.0     0     0   7.9250         3
3             4         1 female  35.0     1     0  53.1000         3
4             5         0   Male  35.0     0     0   8.0500         3

    female          2          3          2          3
0  False  False    True  False    True
1   True  False  False  False  False
2   True  False    True  False    True
3   True  False    True  False  False
4  False  False    True  False    True
```

```
[32]: X = titanic_data.drop("Survived", axis=1)
      Y = titanic_data["Survived"]
```

```
[34]: from sklearn.model_selection import train_test_split
```

```
[39]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3,
      ↪random_state=1)
```

```
[40]: from sklearn.linear_model import LogisticRegression
```

```
[41]: logmodel=LogisticRegression()
```

```
[45]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
      from sklearn.preprocessing import OneHotEncoder

      # Assuming 'X' is your pandas DataFrame and contains categorical columns
      ohe = OneHotEncoder(handle_unknown='ignore') # Create OneHotEncoder instance
      X_encoded = ohe.fit_transform(X.select_dtypes(include=['object'])) # Encode
      ↪categorical columns
```

```

# Convert the encoded data back to a DataFrame for easier handling
X_encoded_df = pd.DataFrame.sparse.from_spmatrix(X_encoded)

# Get feature names for one-hot encoded columns
# Assuming ohe.get_feature_names_out() is available in your sklearn version
feature_names = ohe.get_feature_names_out(X.select_dtypes(include=['object']).
    ↪columns)
X_encoded_df.columns = feature_names # Assign feature names to encoded columns

# Concatenate the encoded categorical features with numerical features if any
X_final = pd.concat([X.select_dtypes(exclude=['object']), X_encoded_df], axis=1)

# Now proceed with splitting and model fitting
X_train, X_test, Y_train, Y_test = train_test_split(X_final, Y, test_size=0.3, ↪
    ↪random_state=1)

# Ensure all column names are strings
X_train.columns = X_train.columns.astype(str)

logmodel = LogisticRegression()
logmodel.fit(X_train, Y_train)

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:768:
UserWarning: pandas.DataFrame with sparse columns found.It will be converted to
a dense numpy array.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

[45]: LogisticRegression()

[46]: predictions = logmodel.predict(X_test)

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:768:
UserWarning: pandas.DataFrame with sparse columns found.It will be converted to
a dense numpy array.
    warnings.warn(

```

```
[47]: from sklearn.metrics import classification_report
```

```
[48]: classification_report(Y_test, predictions)
```

```
[48]: '          precision    recall  f1-score   support\n\n 0.86      0.87      0.86       166\n 101\n\n accuracy          0.83       267\n 0.82      0.81      0.82       267\n\n weighted avg          0.83       0.83       0.83\n 267'
```

```
[49]: from sklearn.metrics import confusion_matrix
```

```
[50]: c = confusion_matrix(Y_test, predictions)
```

```
[52]: confusion_matrix(Y_test, predictions)
```

```
[52]: array([[144,  22],\n        [ 24,  77]])
```

```
[53]: from sklearn.metrics import accuracy_score
```

```
[54]: accuracy_score(Y_test, predictions)
```

```
[54]: 0.8277153558052435
```

```
[ ]:
```

```
[ ]:
```